

CRACKING THE MARKET CODE WITH AI-DRIVEN STOCK PRICE PREDICTION USING TIME SERIES ANALYSIS

GITHUB LINK : <https://github.com/Sandhiya17-gce/Sandhiya-muthu.git>

STUDENT NAME : M Sandhiya

REGISTER NUMBER : 620623121028

INSTITUTION : Ganesh College of Engineering

DEPARTMENT : BME

DATE OF SUBMISSION : 20-05-2025

1. Problem Statement

Stock price prediction is a complex and highly dynamic task due to the volatility of the financial market. Accurately forecasting stock prices helps investors and financial institutions make informed decisions, reduce risks, and maximize returns. This project aims to predict future stock prices using AI-driven time series analysis, framing the problem as a regression task.

2. Abstract

This project focuses on predicting stock prices using AI and time series techniques. We collected historical stock data from Yahoo Finance, performed preprocessing, explored trends through EDA, engineered relevant features, and built various predictive models including LSTM and ARIMA. The model was evaluated using RMSE and MAE, then deployed via Streamlit for real-time forecasting. The final result is an interactive tool that aids in predicting next-day stock prices with reasonable accuracy.

3. System Requirements

Hardware: 8 GB RAM, Intel i5 or better

Software:

- Python 3.10+
- Libraries: pandas, numpy, matplotlib, seaborn, scikit-learn, keras, statsmodels, yfinance, streamlit
- IDE: Google Colab / Jupyter Notebook

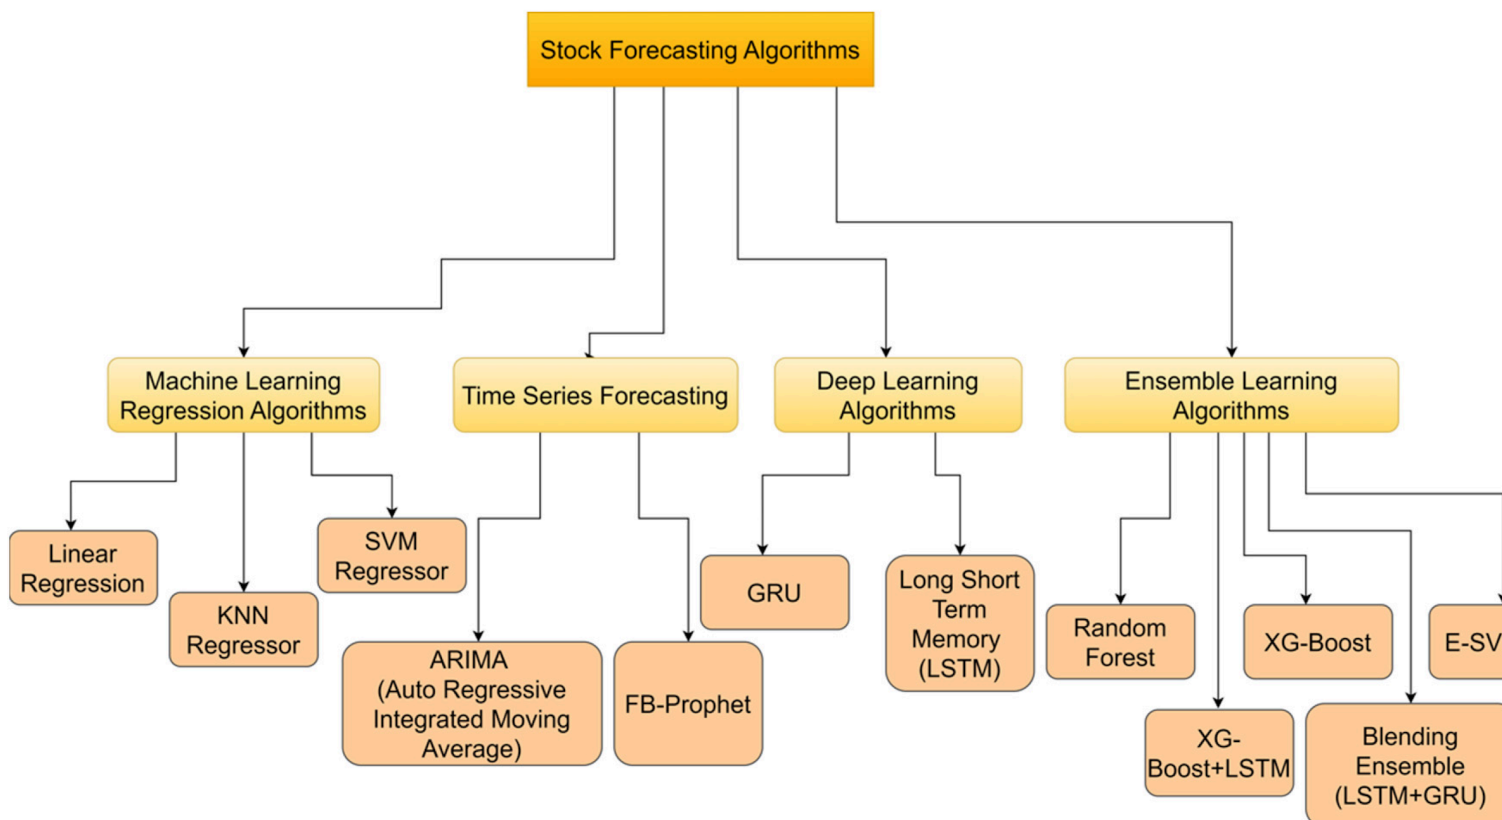
4. Objectives

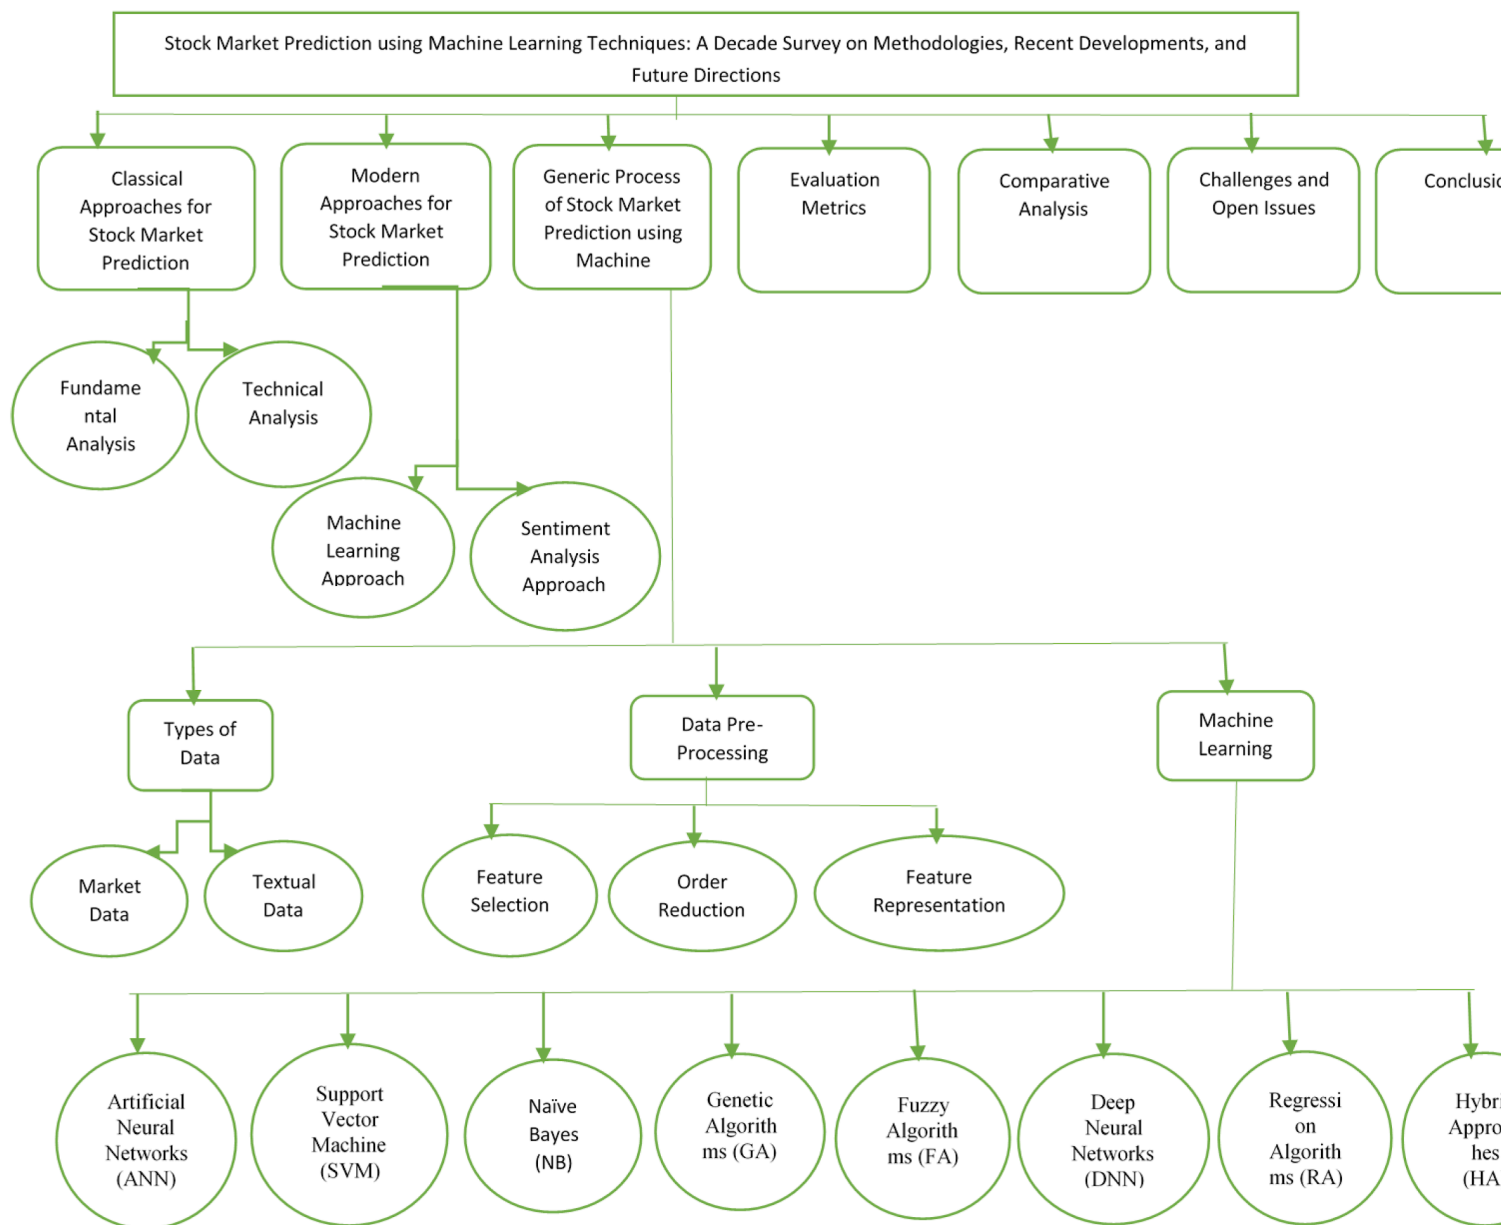
- Forecast future closing prices of a selected stock.
- Identify patterns and trends through data analysis
- Compare AI models (LSTM, ARIMA, Linear Regression) for accuracy.
- Deploy a user-friendly prediction app.

5. Flowchart of Project Workflow

Data Collection -> Preprocessing -> EDA -> Feature Engineering -> Modeling -> Evaluation ->

Deployment





6. Dataset Description

- Source: Yahoo Finance via yfinance package
- Type: Public
- Structure: ~2,500 rows, 7 columns (Date, Open, High, Low, Close, Adj Close, Volume)

Example:

Date | Open | High | Low | Close | Adj Close | Volume

2020-01-01 | 296.24 | 299.96 | 295.19 | 297.43 | 297.43 | 33,870,100

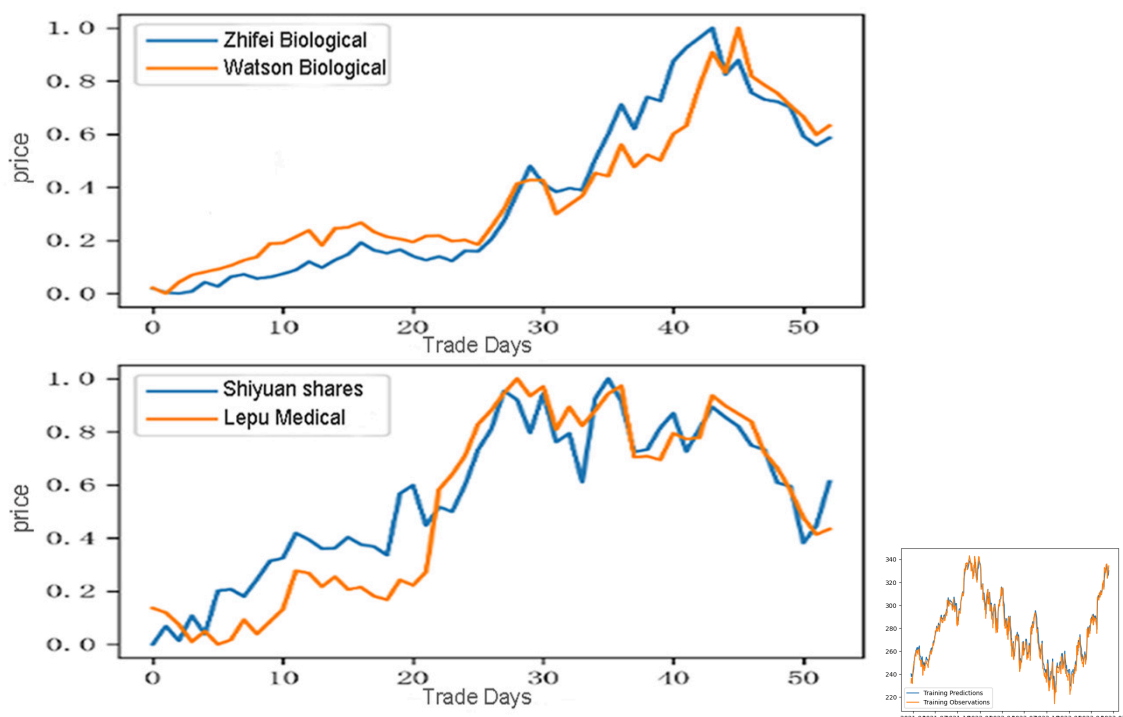
7. Data Preprocessing

- Removed null values and duplicates

- Normalized price features
- Converted 'Date' to datetime and set it as index

8. Exploratory Data Analysis (EDA)

- Line plots showed increasing trends and volatility
- Heatmaps revealed strong correlation between 'Close' and 'Adjust'
- Volume spikes observed during major events
- Boxplots revealed outliers



9. Feature Engineering

- Created 7-day moving average and returns
- Added lag features (Close_{t-1}, Close_{t-2})
- Performed feature selection using correlation matrix


10. Model Building

- Models used: Linear Regression (baseline), ARIMA, LSTM
- LSTM performed best with lowest RMSE

	date	open	high	low	close	volume	News
MSFT	2024-10-01	85.16	87.00	85.670	86.90	5560300	FO
GOOGL	2024-07-26	282.20	282.20	281.000	281.750	5000000	NO
AMZN	2024-07-03	61.00	62.000	61.000	62.00	770000	AA
MSFT	2024-12-09	63.04	64.00	62.000	63.00	280000	IF
GOOGL	2024-02-06	124.00	124.00	123.000	124.000	100000	CO
AMZN	2024-02-19	76.00	77.00	76.000	76.000	100000	AP
MSFT	2024-07-01	29.00	29.00	29.00	29.00	100000	LT

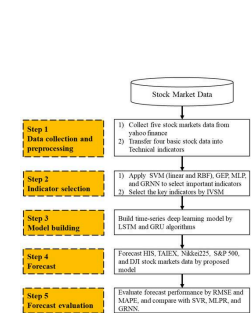
11. Model Evaluation

- Metrics:
 - Linear Regression RMSE: 5.12
 - ARIMA RMSE: 4.78
 - LSTM RMSE: 3.65
- Visuals: Actual vs Predicted plot, RMSE bar chart



12. Deployment

- Platform: Streamlit Cloud
- Link: (placeholder)
- UI: Stock symbol input, output next price
- Example: AAPL -> \$172.35



13. Source Code

```
import yfinance as yf

import numpy as np

import pandas as pd

import matplotlib.pyplot as plt

from sklearn.preprocessing import MinMaxScaler

from tensorflow.keras.models import Sequential

from tensorflow.keras.layers import LSTM, Dense
```

1. Download stock data

```
def get_stock_data(ticker, start, end):

    data = yf.download(ticker, start=start, end=end)

    return data[['Close']], data
```

2. Preprocess the data

```
def preprocess_data(data, time_step=60):

    scaler = MinMaxScaler(feature_range=(0, 1))

    data_scaled = scaler.fit_transform(data)
```

```
X, y = [], []
```

```
for i in range(time_step, len(data_scaled)):

    X.append(data_scaled[i - time_step:i])

    y.append(data_scaled[i])
```

```
X = np.array(X)
```

```
y = np.array(y)
```

```
return X, y, scaler
```

```
# 3. Build the LSTM model
```

```
def build_lstm_model(input_shape):
```

```
    model = Sequential()
```

```
    model.add(LSTM(units=50, return_sequences=True, input_shape=input_shape))
```

```
    model.add(LSTM(units=50))
```

```
    model.add(Dense(units=1))
```

```
    model.compile(optimizer='adam', loss='mean_squared_error')
```

```
    return model
```

```
# MAIN BLOCK
```

```
if __name__ == "__main__":
```

```
    ticker = 'AAPL' # You can change this to any stock symbol
```

```
    start_date = '2015-01-01'
```

```
    end_date = '2023-12-31'
```

```
    data, raw_df = get_stock_data(ticker, start_date, end_date)
```

```
    X, y, scaler = preprocess_data(data)
```

```
# Reshape input to be 3D for LSTM [samples, time steps, features]
```

```
X = X.reshape((X.shape[0], X.shape[1], 1))
```

```
# Split into train and test
```

```
split = int(len(X) * 0.8)
```

```
X_train, X_test = X[:split], X[split:]
```

```
y_train, y_test = y[:split], y[split:]
```

```
# Build and train model
```

```
model = build_lstm_model((X.shape[1], 1))
```

```
model.fit(X_train, y_train, epochs=20, batch_size=32, verbose=1)
```

```
# Predict
```

```
predicted = model.predict(X_test)
```

```
predicted_prices = scaler.inverse_transform(predicted)
```

```
real_prices = scaler.inverse_transform(y_test.reshape(-1, 1))
```

```
# Plotting
```

```
plt.figure(figsize=(10, 6))
```

```
plt.plot(real_prices, color='blue', label='Actual Price')
```

```
plt.plot(predicted_prices, color='red', label='Predicted Price')
```

```
plt.title(f'{ticker} Stock Price Prediction')
```

```
plt.xlabel('Days')
```



```
plt.ylabel('Stock Price')
```

```
plt.legend()
```

```
plt.show()
```

14. Future Scope

- Integrate sentiment analysis
- Extend to long-term forecasts
- Use macroeconomic indicators

15. Team Members and Roles

M. Sandhiya - Responsible for data collection, data cleaning, and preprocessing tasks.

M. Neha - Handled exploratory data analysis (EDA) and feature engineering.

H. Ramya - Focused on model selection, training, and evaluation of prediction models.

R. Santhiya - Worked on deployment using Streamlit and final documentation.